

# **Hyper parameter tuning**

PinkLAB Edu

10 November, 2025

# Table of Contents

1	교차검증.....	4
1.1	교차검증.....	4
1.2	holdout.....	4
1.3	k-fold cross validation .....	4
1.4	stratified k-fold cross validation .....	5
1.5	검증 validation이 끝난 후 test용 데이터로 최종 평가 .....	5
2	교차 검증 구현하기 .....	6
2.1	simple example.....	6
2.2	결과.....	6
2.3	다시 와인 맛 분류하던 데이터로.....	7
2.4	와인 맛 분류기를 위한 데이터 정리 .....	7
2.5	지난번 의사 결정 나무 모델로는?.....	8
2.6	KFold.....	8
2.7	KFold는 index를 반환한다.....	8
2.8	각각의 fold에 대한 학습 후 acc .....	9
2.9	각 acc의 분산이 크지 않다면 평균을 대표 값으로 한다.....	9
2.10	StratifiedKFold .....	10
2.11	acc의 평균이 더 나쁘다.....	10
2.12	cross validation을 보다 간편히 .....	10
2.13	depth가 높다고 무조건 acc가 좋아지는 것도 아니다 .....	11
2.14	train score와 함께 보고 싶다면 .....	11
3	하이퍼파라미터 튜닝 .....	12
3.1	하이퍼파라미터 튜닝 .....	12
3.2	튜닝 대상.....	12
3.3	일단 다시 새파일에서 작업.....	13
3.4	GridSearchCV .....	13
3.5	GridSearchCV의 결과.....	14
3.6	최적의 성능을 가진 모델은?.....	15
3.7	만약 pipeline을 적용한 모델에 GridSearch를 적용하고 싶다면 .....	15

3.8 어렵지 않다 .....	16
3.9 best 모델은? .....	16
3.10 best_score_ .....	16
3.11 Tree 확인해보기 .....	18
3.12 잡기술 하나 - 표로 성능 결과를 정리하자 .....	18

# 1 교차검증


## 1.1 교차검증

- 과적합 : 모델이 학습 데이터에만 과도하게 최적화된 현상.  
그로인해 일반화된 데이터에서는 예측 성능이 과하게 떨어지는 현상
- 지난번 와인 맛 평가에서 훈련용 데이터의 Acc는 72.94,  
테스트용 데이터는 Acc가 71.61%였는데, 누가 이 결과가 정말 괜찮은 것인지 묻는다면?
- 나에게 주어진 데이터에 적용한 모델의 성능을 정확히 표현하기 위해서도 유용하다

## 1.2 holdout




## 1.3 k-fold cross validation




$$\text{Accuracy} = \text{Average}(\text{Accuracy}_1, \dots, \text{Accuracy}_k)$$

## 1.4 stratified k-fold cross validation



$$\text{Accuracy} = \text{Average}(\text{Accuracy}_1, \dots, \text{Accuracy}_k)$$

## 1.5 검증 validation이 끝난 후 test용 데이터로 최종 평가



## 2 교차 검증 구현하기

### 2.1 simple example

```

import numpy as np
from sklearn.model_selection import KFold

X = np.array([[1, 2], [3, 4], [1, 2], [3, 4]])
y = np.array([1, 2, 3, 4])
kf = KFold(n_splits=2)

print(kf.get_n_splits(X))
print(kf)
print('='*30)

for train_idx, test_idx in kf.split(X):
    print("--- idx")
    print(train_idx, test_idx)
    print("--- train data")
    print(X[train_idx])
    print("--- val data")
    print(X[test_idx])
    print('='*30)

```

Python

### 2.2 결과

```

2
KFold(n_splits=2, random_state=None, shuffle=False)
=====
--- idx
[2 3] [0 1]
--- train data
[[1 2]
 [3 4]]
--- val data
[[1 2]
 [3 4]]
=====
--- idx
[0 1] [2 3]
--- train data
[[1 2]
 [3 4]]
--- val data
[[1 2]
 [3 4]]
=====
```

## 2.3 다시 와인 맛 분류하던 데이터로

```
import pandas as pd

red_url = "https://github.com/PinkWink/ML_tutorial/raw/refs/heads" + \
           "/master/dataset/winequality-red.csv"

white_url = "https://github.com/PinkWink/ML_tutorial/raw/refs/heads" + \
           "/master/dataset/winequality-white.csv"

red_wine = pd.read_csv(red_url, sep = ';')
white_wine = pd.read_csv(white_url, sep = ";")

red_wine['color'] = 1.
white_wine['color']= 0.

wine = pd.concat([red_wine, white_wine])
```

✓ 2.0s

Python

## 2.4 와인 맛 분류기를 위한 데이터 정리

```
wine['taste'] = [1. if grade > 5 else 0. for grade in wine['quality']]

X = wine.drop(['taste', 'quality'], axis = 1)
y = wine['taste']
```

✓ 0.0s

Python

## 2.5 지난번 의사 결정 나무 모델로는?

Train Acc : 0.7294593034442948  
Test Acc : 0.7161538461538461

- 여기서 잠깐, 그러니까 누가, “데이터를 저렇게 분리하는 것이 최선인건가?”
  - “저 acc를 어떻게 신뢰할 수 있는가?” 라고 묻는다면~

## 2.6 KFold

```
from sklearn.model_selection import KFold  
  
kfold = KFold(n_splits=5)  
wine_tree_cv = DecisionTreeClassifier(max_depth=2, random_state=13)  
  
0.0s
```

Python

## 2.7 KFold는 index를 반환한다

```
for train_idx, test_idx in kfold.split(X):
    print(len(train_idx), len(test_idx))
0.0s
```

5197 1300  
5197 1300  
5198 1299  
5198 1299  
5198 1299

Python

## 2.8 각각의 fold에 대한 학습 후 acc

```
cv_accuracy = []

for train_idx, test_idx in kfold.split(X, y):
    X_train, X_test = X.iloc[train_idx], X.iloc[test_idx]
    y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]
    wine_tree_cv.fit(X_train, y_train)
    pred = wine_tree_cv.predict(X_test)
    cv_accuracy.append(accuracy_score(y_test, pred))

cv_accuracy
```

✓ 0.0s

```
[0.6007692307692307,
 0.6884615384615385,
 0.7090069284064665,
 0.7628945342571208,
 0.7867590454195535]
```

Python

## 2.9 각 acc의 분산이 크지 않다면 평균을 대표 값으로 한다

```
np.mean(cv_accuracy)

✓ 0.0s
```

np.float64(0.709578255462782)

Python

## 2.10 StratifiedKFold

```
from sklearn.model_selection import StratifiedKFold

skfold = StratifiedKFold(n_splits=5)
wine_tree_cv = DecisionTreeClassifier(max_depth=2, random_state=13)

cv_accuracy = []

for train_idx, test_idx in skfold.split(X, y):
    X_train, X_test = X.iloc[train_idx], X.iloc[test_idx]
    y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]
    wine_tree_cv.fit(X_train, y_train)
    pred = wine_tree_cv.predict(X_test)
    cv_accuracy.append(accuracy_score(y_test, pred))

cv_accuracy
```

✓ 0.0s

[0.5523076923076923,  
 0.6884615384615385,  
 0.7143956889915319,  
 0.7321016166281755,  
 0.7567359507313318]

Python

## 2.11 acc의 평균이 더 나쁘다

```
np.mean(cv_accuracy)
```

✓ 0.0s

np.float64(0.6888004974240539)

Python

- 이런 경우 어떻게 해야 할까?

## 2.12 cross validation을 보다 간편히

```
from sklearn.model_selection import cross_val_score

skfold = StratifiedKFold(n_splits=5)
wine_tree_cv = DecisionTreeClassifier(max_depth=2, random_state=13)

cross_val_score(wine_tree_cv, X, y, scoring=None, cv=skfold)
```

✓ 0.0s

array([0.55230769, 0.68846154, 0.71439569, 0.73210162, 0.75673595])

Python

## 2.13 depth가 높다고 무조건 acc가 좋아지는 것도 아니다

```
wine_tree_cv = DecisionTreeClassifier(max_depth=5, random_state=13)
cross_val_score(wine_tree_cv, X, y, scoring=None, cv=skfold)
✓ 0.0s
array([0.50076923, 0.62615385, 0.69745958, 0.7582756 , 0.74903772])
```

Python

## 2.14 train score와 함께 보고 싶다면


```
from sklearn.model_selection import cross_validate
cross_validate(wine_tree_cv, X, y, scoring=None, cv=skfold,
               return_train_score=True)
✓ 0.0s
{'fit_time': array([0.01190925, 0.01232815, 0.01158357, 0.01125073, 0.01109552]),
 'score_time': array([0.00132251, 0.00120187, 0.00128078, 0.00108957, 0.00108171]),
 'test_score': array([0.50076923, 0.62615385, 0.69745958, 0.7582756 , 0.74903772]),
 'train_score': array([0.78795459, 0.78045026, 0.77568295, 0.76356291,
                      0.76279338])}
```

Python

- 현재 우리는 과적합 현상도 함께 목격하고 있다

## 3 하이퍼파라미터 튜닝

### 3.1 하이퍼파라미터 튜닝



- 모델의 성능을 확보하기 위해 조절하는 설정 값

### 3.2 튜닝 대상

- 결정나무에서 아직 우리가 튜닝해 볼만한 것은 max\_depth이다.
- 간단하게 반복문으로 max\_depth를 바꿔가며 테스트해볼 수 있을 것이다.
- 그런데 앞으로를 생각해서 보다 간편하고 유용한 방법을 생각해보자.

### 3.3 일단 다시 새파일에서 작업

```

import pandas as pd

red_url = (
    "https://github.com/PinkWink/ML_tutorial/raw/refs/heads"
    + "/master/dataset/winequality-red.csv"
)

white_url = (
    "https://github.com/PinkWink/ML_tutorial/raw/refs/heads"
    + "/master/dataset/winequality-white.csv"
)

red_wine = pd.read_csv(red_url, sep=";")
white_wine = pd.read_csv(white_url, sep=";")

red_wine["color"] = 1.0
white_wine["color"] = 0.0

wine = pd.concat([red_wine, white_wine])
wine["taste"] = [1.0 if grade > 5 else 0.0 for grade in wine["quality"]]

X = wine.drop(["taste", "quality"], axis=1)
y = wine["taste"]

```

Python

### 3.4 GridSearchCV

```

from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeClassifier


params = {'max_depth' : [2, 4, 7, 10]}
wine_tree = DecisionTreeClassifier(max_depth = 2, random_state = 13)

gridsearch = GridSearchCV(estimator = wine_tree, param_grid = params , cv = 5)
gridsearch.fit(X, y)

```

✓ 0.3s

Python



- 결과를 확인하고 싶은 파라미터를 정의하면 그만~

- cv는 cross validation

### 3.5 GridSearchCV의 결과

```

import pprint

pp = pprint.PrettyPrinter(indent=4)
pp.pprint(gridsearch.cv_results_)

✓ 0.0s
Python
{
    'mean_fit_time': array([0.00581474, 0.00989447, 0.01649146, 0.022967]),  

    'mean_score_time': array([0.00143118, 0.00139189, 0.00142927, 0.00143967]),  

    'mean_test_score': array([0.6888005, 0.66356523, 0.65340854, 0.64401587]),  

    'param_max_depth': masked_array(data=[2, 4, 7, 10],  

                                      mask=[False, False, False, False],  

                                      fill_value=999999),  

    'params': [ {'max_depth': 2},  

                {'max_depth': 4},  

                {'max_depth': 7},  

                {'max_depth': 10}],  

    'rank_test_score': array([1, 2, 3, 4], dtype=int32),  

    'split0_test_score': array([0.55230769, 0.51230769, 0.50846154, 0.51615385]),  

    'split1_test_score': array([0.68846154, 0.63153846, 0.60307692, 0.60076923]),  

    'split2_test_score': array([0.71439569, 0.72363356, 0.68360277, 0.66743649]),  

    'split3_test_score': array([0.73210162, 0.73210162, 0.73672055, 0.71054657]),  

    'split4_test_score': array([0.75673595, 0.7182448, 0.73518091, 0.72517321]),  

    'std_fit_time': array([0.00051776, 0.00040018, 0.00042915, 0.00078356]),  

    'std_score_time': array([1.93331119e-04, 1.84420370e-04, 1.30381301e-04,  

                           6.80574761e-05]),  

    'std_test_score': array([0.07179934, 0.08390453, 0.08727223, 0.07717557])}

```

### 3.6 최적의 성능을 가진 모델은?

```
gridsearch.best_estimator_
✓ 0.0s Python
```

```
DecisionTreeClassifier(max_depth=2, random_state=13)
```

```
gridsearch.best_score_
✓ 0.0s Python
```

```
np.float64(0.6888004974240539)
```

```
gridsearch.best_params_
✓ 0.0s Python
```

```
{'max_depth': 2}
```

### 3.7 만약 pipeline을 적용한 모델에 GridSearch를 적용하고 싶다면

```
from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import StandardScaler


estimators = [('scaler', StandardScaler()),
              ('clf', DecisionTreeClassifier(random_state = 13))]
pipe = Pipeline(estimators)
✓ 0.0s Python
```

### 3.8 어렵지 않다


```
param_grid = [ {'clf__max_depth' : [2, 4, 7, 10]}]

GridSearch = GridSearchCV(estimator = pipe, param_grid = param_grid, cv = 5)
GridSearch.fit(X, y)

✓ 0.4s
```



### 3.9 best 모델은?



### 3.10 best\_score\_

```
GridSearch.best_score_

✓ 0.0s
```

np.float64(0.6888004974240539)

## GridSearch.cv\_results\_

✓ 0.0s

Python

```
{'mean_fit_time': array([0.00800552, 0.01074991, 0.01913009, 0.02879853]),  
 'std_fit_time': array([0.00196962, 0.000345 , 0.00237366, 0.00145621]),  
 'mean_score_time': array([0.00164223, 0.00155172, 0.00184488, 0.00205579]),  
 'std_score_time': array([0.00029029, 0.00016769, 0.00039364, 0.0003351 ]),  
 'param_clf_max_depth': masked_array(data=[2, 4, 7, 10],  
                                       mask=[False, False, False, False],  
                                       fill_value=999999),  
 'params': [{ 'clf_max_depth': 2},  
            { 'clf_max_depth': 4},  
            { 'clf_max_depth': 7},  
            { 'clf_max_depth': 10}],  
 'split0_test_score': array([0.55230769, 0.51230769, 0.50846154, 0.51615385]),  
 'split1_test_score': array([0.68846154, 0.63153846, 0.60461538, 0.60230769]),  
 'split2_test_score': array([0.71439569, 0.72363356, 0.68206313, 0.66589684]),  
 'split3_test_score': array([0.73210162, 0.73210162, 0.73672055, 0.71054657]),  
 'split4_test_score': array([0.75673595, 0.7182448 , 0.73518091, 0.72517321]),  
 'mean_test_score': array([0.6888005 , 0.66356523, 0.6534083 , 0.64401563]),  
 'std_test_score': array([0.07179934, 0.08390453, 0.08699322, 0.0769154 ]),  
 'rank_test_score': array([1, 2, 3, 4], dtype=int32)}
```


### 3.11 Tree 확인해보기

```
import matplotlib.pyplot as plt
from sklearn import tree

fig = plt.figure(figsize=(15, 8))
= tree.plot_tree(GridSearch.best_estimator_['clf'],
                 feature_names=X.columns,
                 class_names = ["W" , "R"],
                 rounded = True,
                 filled=True)
```

✓ 0.1s

Python



### 3.12 잡기술 하나 - 표로 성능 결과를 정리하자

```
import pandas as pd

score_df = pd.DataFrame(GridSearch.cv_results_)
score_df[['params' , 'rank_test_score', 'mean_test_score' , 'std_test_score']]
```

✓ 0.0s

Python

	params	rank_test_score	mean_test_score	std_test_score
0	{'clf__max_depth': 2}	1	0.688800	0.071799
1	{'clf__max_depth': 4}	2	0.663565	0.083905
2	{'clf__max_depth': 7}	3	0.653408	0.086993
3	{'clf__max_depth': 10}	4	0.644016	0.076915

- accuracy의 평균과 표준편차를 보자~