

DeepLearning with MNIST

PinkLAB Edu

20 November, 2025

Table of Contents

1	MNIST 데이터.....	4
1.1	MNIST 데이터로 한 번 더 DL 과정을 복습해보자.....	4
1.2	Tensorflow에서 MNSIT 읽기	4
1.3	one-hot-encoding	4
1.4	이번 모델~	5
1.5	모델을 만들어보자.....	5
1.6	다시 한 번 더 softmax란?	6
1.7	model.summary	6
1.8	코드와 모델은 일치한다.....	7
1.9	fit~.....	7
1.10	acc와 loss를 그려보자.....	8
1.11	머신러닝에서 93%쯤 나왔던 결과대비 5%쯤 향상되었다	9
1.12	뭐가 틀렸나 확인해보자.....	9
1.13	틀린 데이터의 인덱스만 모아서.....	9
1.14	그중 16개만	10
1.15	뭘 틀렸나?.....	10
1.16	나도 틀릴만한게 있네	10
2	MNIST fashion	11
2.1	MNIST fashion data	11
2.2	데이터 읽기	11
2.3	어떻게 생겼는지 확인해보자	12
2.4	패션~	12
2.5	모델은 숫자때와 동일한 구조로 두자.....	13
2.6	summary	13
2.7	fit~~	14
2.8	학습 상황을 관찰해 보자	14
2.9	잘 학습이 되는듯 하지만,	15
2.10	테스트데이터 accuracy	15
2.11	어떤 데이터가 틀렸는지 추출하고	16

2.12 틀린 데이터를 모아서~	16
2.13 16개만 선택해서.....	16
2.14 그려보자.....	17
2.15 결과.....	17

1 MNIST 데이터

1.1 MNIST 데이터로 한 번 더 DL 과정을 복습해보자

- MNIST 데이터에 대해서는 kNN에서 설명을 했었다.

1.2 Tensorflow에서 MNSIT 읽기

```
import tensorflow as tf

mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
```


Python

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 [=====] - 0s 0us/step

1.3 one-hot-encoding

- 이 타이밍에 one-hot-encoding을 해야 한다.
- 그런데 또 하나의 방법이 loss 함수를 `sparse_categorical_crossentropy`로 설정하면 같은 효과이다.
- 그래서 pass~

1.4 이번 모델~



1.5 모델을 만들어보자


```
model = tf.keras.models.Sequential(
    [
        tf.keras.layers.Flatten(input_shape=(28, 28)),
        tf.keras.layers.Dense(1000, activation="relu"),
        tf.keras.layers.Dense(10, activation="softmax"),
    ]
)

model.compile(
    optimizer="adam", loss="sparse_categorical_crossentropy", metrics=["accuracy"]
)
```

Python

1.6 다시 한 번 더 softmax란?

Multi-Class Classification with NN and SoftMax Function



1.7 model.summary

```
model.summary()
```

✓ 0.0s

Python

Model: "sequential"


Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 1000)	785,000
dense_1 (Dense)	(None, 10)	10,010

Total params: 795,010 (3.03 MB)

Trainable params: 795,010 (3.03 MB)

Non-trainable params: 0 (0.00 B)

1.8 코드와 모델은 일치한다



1.9 fit~

```

import time

start_time = time.time()
hist = model.fit(
    x_train,
    y_train,
    validation_data=(x_test, y_test),
    epochs=10,
    batch_size=100,
    verbose=1,
)
print("Fit time :", time.time() - start_time)
✓ 28.0s
Python

```

600/600 ————— 3s 5ms/step - accuracy: 0.9961 - loss: 0.0132 -
val_accuracy: 0.9812 - val_loss: 0.0672
Epoch 8/10
600/600 ————— 3s 5ms/step - accuracy: 0.9968 - loss: 0.0108 -
val_accuracy: 0.9812 - val_loss: 0.0662
Epoch 9/10
600/600 ————— 3s 5ms/step - accuracy: 0.9968 - loss: 0.0099 -
val_accuracy: 0.9803 - val_loss: 0.0729
Epoch 10/10
600/600 ————— 3s 5ms/step - accuracy: 0.9976 - loss: 0.0076 -
val_accuracy: 0.9801 - val_loss: 0.0708
Fit time : 28.048307180404663

1.10 acc와 loss를 그려보자

```
import matplotlib.pyplot as plt

plot_target = ["loss", "val_loss", "accuracy", "val_accuracy"]


plt.figure(figsize=(12, 8))

for each in plot_target:
    plt.plot(hist.history[each], label=each)

plt.legend()
plt.grid()
plt.show()
```

✓ 0.1s

Python



1.11 머신러닝에서 93%쯤 나왔던 결과대비 5%쯤 향상되었다

```

score = model.evaluate(x_test, y_test)
print("Test loss :", score[0])
print("Test accuracy :", score[1])
✓ 0.4s
313/313 ━━━━━━━━ 0s 1ms/step - accuracy: 0.9757 - loss: 0.0866
Test loss : 0.07075094431638718
Test accuracy : 0.9800999760627747

```

Python

1.12 뭐가 틀렸나 확인해보자

```

import numpy as np

predicted_result = model.predict(x_test)
predicted_labels = np.argmax(predicted_result, axis=1)
predicted_labels[:10]
✓ 0.4s
313/313 ━━━━━━━━ 0s 920us/step
array([7, 2, 1, 0, 4, 1, 4, 9, 5, 9])

```

Python

```

y_test[:10]
✓ 0.0s
array([7, 2, 1, 0, 4, 1, 4, 9, 5, 9], dtype=uint8)

```

Python

1.13 틀린 데이터의 인덱스만 모아서

```

wrong_result = []

for n in range(0, len(y_test)):
    if predicted_labels[n] != y_test[n]:
        wrong_result.append(n)

len(wrong_result)
✓ 0.0s

```

Python

199

1.14 그중 16개만

```
import random

samples = random.choices(population=wrong_result, k=16)
samples
```

✓ 0.0s

Python

[7216,
 9944,
 3073,
 8504,
 720,
 9634,
 965,
 1247,
 2730,
 3941,

1.15 뭘 틀렸나?

```
plt.figure(figsize=(10, 10))

for idx, n in enumerate(samples):
    plt.subplot(4, 4, idx + 1)
    plt.imshow(x_test[n].reshape(28, 28), cmap="Greys")
    plt.title("Label : " + str(y_test[n]) + " | Predict : " + str(predicted_labels[n]))
    plt.axis("off")

plt.show()
```


✓ 0.4s

Python

1.16 나도 틀릴만한게 있네

2 MNIST fashion

2.1 MNIST fashion data



- 숫자로 된 MNIST데이터처럼 28*28 크기의 패션과 관련된 10개 종류의 데이터

2.2 데이터 읽기

```
import tensorflow as tf

fashion_mnist = tf.keras.datasets.fashion_mnist

(X_train, y_train), (X_test, y_test) = fashion_mnist.load_data()
X_train, X_test = X_train / 255.0, X_test / 255.0
✓ 2.7s
```

Python

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>
29515/29515 **0s** 0us/step
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>
26421880/26421880 **1s** 0us/step
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>
5148/5148 **0s** 0us/step
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>
4422102/4422102 **0s** 0us/step

2.3 어떻게 생겼는지 확인해보자

```
import random
import matplotlib.pyplot as plt

samples = random.choices(population=range(0, len(y_train)), k=16)
```

✓ 0.0s

Python

```
plt.figure(figsize=(10, 8))

for idx, n in enumerate(samples):
    plt.subplot(4, 4, idx + 1)
    plt.imshow(X_train[n].reshape(28, 28), cmap="Greys")
    plt.title("Label : " + str(y_train[n]))
    plt.axis("off")

plt.show()
```

✓ 0.4s

Python

2.4 패션~



2.5 모델은 숫자때와 동일한 구조로 두자

```
model = tf.keras.models.Sequential(
    [
        tf.keras.layers.Flatten(input_shape=(28, 28)),
        tf.keras.layers.Dense(1000, activation="relu"),
        tf.keras.layers.Dense(10, activation="softmax"),
    ]
)

model.compile(
    optimizer="adam", loss="sparse_categorical_crossentropy", metrics=["accuracy"]
)
```

✓ 0.0s Python

2.6 summary

```
model.summary()
```

✓ 0.0s Python

Model: "sequential_1"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 784)	0
dense_2 (Dense)	(None, 1000)	785,000
dense_3 (Dense)	(None, 10)	10,010

Total params: 795,010 (3.03 MB)

Trainable params: 795,010 (3.03 MB)

Non-trainable params: 0 (0.00 B)

2.7 fit~~

```

import time

start_time = time.time()
hist = model.fit(
    x_train,
    y_train,
    validation_data=(x_test, y_test),
    epochs=10,
    batch_size=100,
    verbose=1,
)
print("Fit time :", time.time() - start_time)
✓ 34.4s

```

Python

```

Epoch 6/10
600/600 ━━━━━━━━ 3s 6ms/step - accuracy: 0.1499 - loss: 2.2588 -
val_accuracy: 0.1012 - val_loss: 2.3372
Epoch 7/10
600/600 ━━━━━━ 3s 6ms/step - accuracy: 0.1704 - loss: 2.2330 -
val_accuracy: 0.0995 - val_loss: 2.3507
Epoch 8/10
600/600 ━━━━ 3s 5ms/step - accuracy: 0.1885 - loss: 2.2024 -
val_accuracy: 0.1017 - val_loss: 2.3678
Epoch 9/10
600/600 ━━━━ 3s 5ms/step - accuracy: 0.2091 - loss: 2.1627 -
val_accuracy: 0.0939 - val_loss: 2.3917
Epoch 10/10
600/600 ━━━━ 3s 6ms/step - accuracy: 0.2319 - loss: 2.1176 -
val_accuracy: 0.0990 - val_loss: 2.4227
Fit time : 34.417144775390625

```

2.8 학습 상황을 관찰해 보자

```

import matplotlib.pyplot as plt

plot_target = ["loss", "val_loss", "accuracy", "val_accuracy"]

plt.figure(figsize=(12, 8))


for each in plot_target:
    plt.plot(hist.history[each], label=each)

plt.legend()
plt.grid()
plt.show()
✓ 0.1s

```

Python

2.9 잘 학습이 되는듯 하지만,



- val_loss와 train loss사이에 간격이 발생한다.

2.10 테스트데이터 accuracy

```
score = model.evaluate(X_test, y_test)
print("Test loss :", score[0])
print("Test accuracy :", score[1])
```

✓ 0.6s

Python

```
313/313 ━━━━━━━━ 1s 1ms/step - accuracy: 0.8747 - loss: 0.3500
Test loss : 0.34942078590393066
Test accuracy : 0.8751999735832214
```

2.11 어떤 데이터가 틀렸는지 추출하고

```
import numpy as np

predicted_result = model.predict(X_test)
predicted_labels = np.argmax(predicted_result, axis=1)
predicted_labels[:10]

✓ 0.5s
```

313/313 ————— 0s 1ms/step

```
array([9, 2, 1, 1, 0, 1, 4, 6, 5, 7])
```

Python

```
y_test[:10]

✓ 0.0s
```

```
array([9, 2, 1, 1, 6, 1, 4, 6, 5, 7], dtype=uint8)
```

Python

2.12 틀린 데이터를 모아서~

```
wrong_result = []

for n in range(0, len(y_test)):
    if predicted_labels[n] != y_test[n]:
        wrong_result.append(n)

len(wrong_result)

✓ 0.0s
```

Python

2.13 16개만 선택해서

```
import random

samples = random.choices(population=wrong_result, k=16)
samples

✓ 0.0s
```

```
[5052,
 7853,
 4159,
 9126,
 1396,
 316,
 3084,
 344,
 2507,
 ...]
```

Python

2.14 그려보자

```
plt.figure(figsize=(10, 8))

for idx, n in enumerate(samples):
    plt.subplot(4, 4, idx + 1)
    plt.imshow(X_train[n].reshape(28, 28), cmap="Greys", interpolation="nearest")
    plt.title("Label : " + str(y_train[n]) + " Predict : " + str(predicted_labels[n]))
    plt.axis("off")

plt.show()

✓ 0.2s
```

Python

2.15 결과



- 0 : 티셔츠
- 1 : 바지
- 2 : 스웨터

- 3 : 드레스
- 4 : 코트
- 5 : 샌들
- 6 : 셔츠
- 7 : 운동화
- 8 : 가방
- 9 : 부츠